

---

Note: For easy navigation, please use the Table of contents

## **Section 0: Name**

Pamela Sin Hui

## **Section 1: Project Title**

Product Detection Image Classification - Using Kerastuner, Transfer Learning, Finetuning

## ▼ Section 2: Goals, Dataset, Tasks





### 1. Goals

In Shopee Product Detection Dataset, there are more than 100k images directly from E-commerce industry field. You will be able to explore the real-world images which is noisy and long-tailed, and let your model predict the correct categories for the images. There contains 42 most popular categories product at Shopee.

<https://www.kaggle.com/c/shopee-product-detection-open/overview>

### 2. Dataset

- **train folder** - Contains 42 folders (each labelled '00', '01', ..., '41'), each containing images of 42 popular product categories at Shoppee.
  - **test folder** - Contains images where the goal is to give labels (from the 42 categories) to each image.
  - **train.csv** - 2 Columns (filename and category)
  - **test.csv** - 2 Columns (filename and dummy category)
- 
- 

### 3. Tasks

Note: We re-write certain blocks of code again in some cells due to the possibility of Colab-run timeouts

- Data exploration
- Create train and validation datasets from smaller sample subset of 3000 images - For more rapid coding and testing
- Transfer learning (mobilenet and inception v3 featurizers) + Kerastuner randomsearch
- Transfer learning + Hyperparameter search with tensorboard visualisation
- Perform further training on best Transfer Learning featurizer + hyperparameters
- Check model performance and metrics
- Train and validate on global dataset
- Ideas for improvement

```

1 import datetime, os
2 import shutil
3 import zipfile
4 import pickle
5
6 import numpy as np
7 import pandas as pd
8 import matplotlib.pyplot as plt
9
10 from sklearn.metrics import classification_report
11
12 import tensorflow as tf
13 from keras.layers import Dropout
14 from tensorflow.keras.applications import MobileNetV2
15 from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
16 from tensorflow.keras.applications.inception_v3 import InceptionV3
17 from tensorflow.keras.preprocessing.image import ImageDataGenerator
18 from tensorflow.keras import layers
19 from tensorflow.keras.layers import Input, Conv2D, LeakyReLU, Flatten, Dense,
20 from tensorflow.keras import Model
21 from tensorflow.keras.callbacks import ModelCheckpoint
22 from tensorflow.keras.models import load_model
23 from tensorflow.keras.optimizers import RMSprop
24 from tensorflow.keras.optimizers import SGD

```

```

1 #DO NOT RUN THIS CELL AFTER FIRST RUN
2 #Uploading the zip file to Google Drive took ~3 hours
3 #The below unzip operation took ~ 45 min, but did not unzip folders 15 to 41
4 #Manual unzip locally was done ~45 min, and manual drag and drop upload of fo
5 !unzip "/content/drive/My Drive/shopee-product-detection-open.zip" -d "/conte

```

#### Streaming output truncated to the last 5000 lines.

```

inflating: /content/drive/My Drive/train/train/train/13/31c2dc64a0007d1f3
inflating: /content/drive/My Drive/train/train/train/13/31c9a82cf3f38dcfd
inflating: /content/drive/My Drive/train/train/train/13/31dc15cb2eb3f2c91
inflating: /content/drive/My Drive/train/train/train/13/31e4d096df72eaaca
inflating: /content/drive/My Drive/train/train/train/13/31ebeaaa50685c435
inflating: /content/drive/My Drive/train/train/train/13/320ab3dfa81199c85
inflating: /content/drive/My Drive/train/train/train/13/32174db1e4b57b5c4
inflating: /content/drive/My Drive/train/train/train/13/32209e1ac67fb78f7
inflating: /content/drive/My Drive/train/train/train/13/322143750fb1e4717
inflating: /content/drive/My Drive/train/train/train/13/322df0ed3b7ffe60f
inflating: /content/drive/My Drive/train/train/train/13/3230da734295bebf5
inflating: /content/drive/My Drive/train/train/train/13/3232afc5ae1762af9
inflating: /content/drive/My Drive/train/train/train/13/325572f3ffc7950da
inflating: /content/drive/My Drive/train/train/train/13/325a78e4243ec15ee
inflating: /content/drive/My Drive/train/train/train/13/3267f4caf8bc4823f
inflating: /content/drive/My Drive/train/train/train/13/327250586f75d4ba5
inflating: /content/drive/My Drive/train/train/train/13/32d22a2cdc1b0eb15
inflating: /content/drive/My Drive/train/train/train/13/3364558dda36d0ab0
inflating: /content/drive/My Drive/train/train/train/13/336fffd7e5eab0feb
inflating: /content/drive/My Drive/train/train/train/13/337fac602c56672ff
inflating: /content/drive/My Drive/train/train/train/13/338a44430a849d446

```

inflating: /content/drive/My Drive/train/train/train/13/33bb74713f53595ff  
inflating: /content/drive/My Drive/train/train/train/13/33c71d4f438084112  
inflating: /content/drive/My Drive/train/train/train/13/33d3862ec9c022e60  
inflating: /content/drive/My Drive/train/train/train/13/34217c80a8c4a7dcc  
inflating: /content/drive/My Drive/train/train/train/13/342de12b5f9f4f80c  
inflating: /content/drive/My Drive/train/train/train/13/3447112e0247f4830  
inflating: /content/drive/My Drive/train/train/train/13/34c42328e2389febcb  
inflating: /content/drive/My Drive/train/train/train/13/34effad42d6fc0e4e  
inflating: /content/drive/My Drive/train/train/train/13/34f2cf349431d551c  
inflating: /content/drive/My Drive/train/train/train/13/351b96269366a3d93  
inflating: /content/drive/My Drive/train/train/train/13/35212edd9057a9a0f  
inflating: /content/drive/My Drive/train/train/train/13/3525f1758798a6884  
inflating: /content/drive/My Drive/train/train/train/13/3526ba6a256d1f159  
inflating: /content/drive/My Drive/train/train/train/13/3548a1b532690c665  
inflating: /content/drive/My Drive/train/train/train/13/359fc4987220dc8e9  
inflating: /content/drive/My Drive/train/train/train/13/35a4530e0d2339373  
inflating: /content/drive/My Drive/train/train/train/13/364b1db2c553d4868  
inflating: /content/drive/My Drive/train/train/train/13/3684088dd79bd46a0  
inflating: /content/drive/My Drive/train/train/train/13/36985e9089e0d7ca4  
inflating: /content/drive/My Drive/train/train/train/13/36fc3178b3dffb0b4e  
inflating: /content/drive/My Drive/train/train/train/13/37063530fb8b4559c  
inflating: /content/drive/My Drive/train/train/train/13/370b537f474a48d78  
inflating: /content/drive/My Drive/train/train/train/13/3714fb099b5dcbfab  
inflating: /content/drive/My Drive/train/train/train/13/37452f745da28286e  
inflating: /content/drive/My Drive/train/train/train/13/3756c1aa3374bdf37  
inflating: /content/drive/My Drive/train/train/train/13/37614f5f974d4bb8b  
inflating: /content/drive/My Drive/train/train/train/13/377db5d2f8bf82e29  
inflating: /content/drive/My Drive/train/train/train/13/37850a979ca9bf98b  
inflating: /content/drive/My Drive/train/train/train/13/37ca866e2e6bbc079  
inflating: /content/drive/My Drive/train/train/train/13/37cb389558ba72529  
inflating: /content/drive/My Drive/train/train/train/13/37d1afc9e17fa4b36  
inflating: /content/drive/My Drive/train/train/train/13/37e320be999fcffbd  
inflating: /content/drive/My Drive/train/train/train/13/381eaf2910c4fb54a  
inflating: /content/drive/My Drive/train/train/train/13/383c62550c4f6eca8  
inflating: /content/drive/My Drive/train/train/train/13/387a3a90749478e3c  
inflating: /content/drive/My Drive/train/train/train/13/387ca76cb70a51509  
inflating: /content/drive/My Drive/train/train/train/13/389558be25a294f72  
inflating: /content/drive/My Drive/train/train/train/13/38a288dc1d27c24b

## ▼ Section 3: Data Engineering

- Data exploration
- Prepare smaller sample subset of 3000 images

```
1 df_traincsv = pd.read_csv('/content/drive/My Drive/train.csv')
2 df_traincsv.head()
```

	filename	category
0	45e2d0c97f7bdf8cbf3594beb6fdcda0.jpg	3
1	f74d1a5fc2498bbbfa045c74e3cc333e.jpg	3
2	f6c172096818c5fab10ecae722840798.jpg	3
3	251ffd610399ac00fea7709c642676ee.jpg	3
4	73c7328b8eda399199fdedec6e4badaf.jpg	3

```
1 df_traincsv.shape

(105390, 2)
```

```
1 df_traincsv.dtypes

filename    object
category    int64
dtype: object
```

```
1 df_check_files = df_traincsv.groupby(['category']).count()
2 df_check_files.head()
```

	filename
category	
0	2683
1	2702
2	2687
3	2703
4	2703

```
1 #Sanity check - Count number of files in the original train dataset path upto
2 import os
3 df_check_files['directory_count'] = 0
4 for i in range(42):
5     img_folder_path = "/content/drive/My Drive/train/train/train/" + folder_name
6     dirListing = os.listdir(img_folder_path)
7     df_check_files.loc[i,'directory_count'] = len(dirListing)
```

```

1 #The steel tower was a 100K images in original and then dataset comparison of th
2 df_check_files['difference'] = df_check_files.filename - df_check_files.direc
3 df_check_files

```

	filename	directory_count	difference
category			
0	2683	2281	402
1	2702	2298	404
2	2687	2284	403
3	2703	2298	405
4	2703	2298	405
5	2641	2245	396
6	2641	2245	396
7	2660	2262	398
8	2700	2295	405
9	2698	2293	405
10	2672	2272	400
11	1843	1567	276
12	2691	2287	404
13	2682	2280	402
14	2684	2282	402
15	2632	2237	395
16	2665	2265	400
17	1553	1320	233
18	2103	1789	314
19	2679	2337	342
20	2653	2256	397
21	2598	2208	390
22	2623	2230	393
23	2540	2159	381
24	2705	2299	406
25	2692	2288	404
26	2684	2281	403

27	2702	2297	405
28	2561	2177	384
29	2138	1817	321
30	1702	1400	250

```

1 #Creating a list of folder names from '00' to '41'
2 folder_names = []
3 for i in range(0,42,1):
4     if len(str(i)) == 1:
5         folder_names.append(str(0) + str(i))
6     else:
7         folder_names.append(str(i))
8 print(folder_names)

```

```
['00', '01', '02', '03', '04', '05', '06', '07', '08', '09', '10', '11', '12', '13', '14', '15', '16', '17', '18', '19', '20', '21', '22', '23', '24', '25', '26', '27', '28', '29', '30', '31', '32', '33', '34', '35', '36', '37', '38', '39', '40', '41']
```

```

1 #Create new_train_df dataframe by checking which image filenames are in each folder
2 #We had split the data into train and validation sets, so we can get the file names
3 import os
4 from os import listdir
5 from os.path import isfile, join
6 the_index = 0
7 new_train_df = pd.DataFrame(columns = ['filename', 'category'])
8 for i in range(42):
9     img_folder_path = "/content/drive/My Drive/train/train/train/" + folder_names[i]
10    onlyfiles = [f for f in listdir(img_folder_path) if isfile(join(img_folder_path, f))]
11    temp_df = pd.DataFrame(columns = ['filename', 'category'])
12    temp_df.filename = onlyfiles
13    temp_df.category = i
14    new_train_df = new_train_df.append(temp_df, ignore_index=True)

```

```
1 new_train_df.head()
```

	filename	category
0	d089953699c05da67da914cceb991d5e.jpg	0
1	d0018d2c844645f14c0e33de7415348d.jpg	0
2	d2178eab8cb870b65a396cd1808f6528.jpg	0
3	c6187c65a9883ffae35c9e3301a32417.jpg	0
4	ce3a9f520751cc48de872d500ecd541a.jpg	0

```
1 new_train_df.shape
```

```
(89646, 2)
```

```

1 #Show random 3 images from each category
2 import matplotlib.image as mpimg
3 for i in range(42):
4     fig=plt.figure(figsize=(5, 5))
5     img_folder_path = "/content/drive/My Drive/train/train/train/" + folder_name
6     temp_list = list(new_train_df.loc[new_train_df.category == i].sample(n=3).files)
7     print('Category ' + str(i) + ': ')
8     for j in range(3):
9         img = mpimg.imread(img_folder_path + '/' + temp_list[j])
10        fig.add_subplot(1, 3, j+1)
11        plt.axis('off')
12        plt.title(str(i))
13        plt.imshow(img)
14    plt.show()

```

Category 0:



Category 1:



Category 2:




Category 3:



Category 4:



Category 5:







Category 6:

6



6



6



Category 7:

7



7



7



```
1 df_traincsv.head()
```

	filename	category
0	45e2d0c97f7bdf8cbf3594beb6fdcdca0.jpg	3
1	f74d1a5fc2498bbbfa045c74e3cc333e.jpg	3
2	f6c172096818c5fab10ecae722840798.jpg	3
3	251ffd610399ac00fea7709c642676ee.jpg	3
4	73c7328b8eda399199fdedec6e4badaf.jpg	3



```
1 df_testcsv = pd.read_csv('/content/drive/My Drive/test.csv')
2 df_testcsv.head()
```

	filename	category
0	fd663cf2b6e1d7b02938c6aaaae0a32d2.jpg	43
1	c7fd77508a8c355eaab0d4e10efd6b15.jpg	43
2	127f3e6d6e3491b2459812353f33a913.jpg	43
3	5ca4f2da11eda083064e6c36f37eeb81.jpg	43
4	46d681a542f2c71be017eef6aae23313.jpg	43



```
1 df_testcsv.shape
```

(12186, 2)



```

1 img_folder_path = "/content/drive/My Drive/test/test/test"
2 dirListing = os.listdir(img_folder_path)
3 len(dirListing)

```

12192



Can consider making this sampling logic into a function and adjust N=3000

```

1 #DO NOT RUN THIS CELL AFTER FIRST RUN
2 #We will take a sample of 3000 images from the train dataset to build initial
3 #From the sample of 3000 images, we will also split into train and validation
4 #We create a df_traincsv_3000marker dataframe to later note which photos have
5 df_traincsv_3000marker = df_traincsv
6 df_traincsv_3000marker['selected_for_3000'] = 0
7 df_traincsv_3000marker['selected_for_3000_validation'] = 0
8 df_traincsv_3000marker.head()

```

	filename	category	selected_for_3000	selected_for_3000_validation
0	45e2d0c97f7bdf8cbf3594beb6fdcdca0.jpg	3	0	0
1	f74d1a5fc2498bbbfa045c74e3cc333e.jpg	3	0	0
2	f6c172096818c5fab10ecae722840798.jpg	3	0	0
3	251ffd610399ac00fea7709c642676ee.jpg	3	0	0
4	73c7328b8eda399199fddedec6e4badaf.jpg	3	0	0



```

1 #DO NOT RUN THIS CELL AFTER FIRST RUN
2 #Taking a sample of 3000 from the train dataset
3 #Creating folders '00' to '41', and sampling 3% of each category
4 #Images are COPIED, leaving the original test dataset as is
5 rolling = 0
6 for two_digit in folder_names:
7     path = "/content/drive/My Drive/sample_3000/train/" + two_digit
8     os.mkdir(path)
9
10    temp_df = df_traincsv.loc[df_traincsv.category == rolling,:].sample(frac=0.03)
11    temp_index = temp_df.index
12    df_traincsv_3000marker.loc[temp_index, 'selected_for_3000'] = 1
13    temp_df.reset_index(inplace=True)
14
15    for i in range(len(temp_df)):
16        f = temp_df.loc[i, 'filename']
17        shutil.copy("/content/drive/My Drive/train/train/train/" + two_digit + '/' + f, path)
18
19    rolling+=1

```




Category 20:



```

1 #DO NOT RUN THIS CELL AFTER FIRST RUN
2 #TO BE REPEATED (WITH SOME CHANGE) DURING RUN ON MORE IMAGES
3 #From the sample of 3000, we MOVE 20% of images in each category to the valid
4 rolling = 0
5 for two_digit in folder_names:
6     path = "/content/drive/My Drive/sample_3000/validation/" + two_digit
7     os.mkdir(path)
8
9     df_traincsv_3000marker
10    temp_df = df_traincsv_3000marker.loc[(df_traincsv_3000marker.category==roll
11    temp_index = temp_df.index
12    df_traincsv_3000marker.loc[temp_index, 'selected_for_3000_validation'] = 1
13    temp_df.reset_index(inplace=True)
14
15    for i in range(len(temp_df)):
16        f = temp_df.loc[i, 'filename']
17        shutil.move("/content/drive/My Drive/sample_3000/train/" + two_digit + '/'
18
19    rolling+=1

```



## Section 4: Feature Engineering

- Use ImageDataGenerator to generate augmented images as the sample 3000 dataset is small



## ▼ Section 5: Model Engineering

- Transfer learning - Featurizers from mobilenet and inceptionv3 pre-trained CNNs
- Kerastuner RandomSearch - No. of hidden layers, no. units in hidden layers, Optimizer
- Hyperparameter search with tensorboard visualisation - No. units in hidden layer1, no. units in hidden layer2, Optimizer, % dropout



```

1 featurizer_mobilenet = MobileNetV2(input_shape=(224, 224, 3), include_top=False)
2 featurizer_inception = InceptionV3(input_shape = (299, 299, 3), include_top = False)
3 #featurizer_mobilenet.summary()
4 #featurizer_inception.summary()

```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/9412608/9406464> [=====] - 0s 0us/step  
 Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/87916544/87910968> [=====] - 1s 0us/step



## ▼ --- Kerastuner - RandomSearch ---



```

1 #mobilenet featurizer and kerastuner RandomSearch on 1. Hidden layers (2 or 3
2 #Running for just 2 iterations – We will later use searching across another s
3 !pip install keras-tuner
4 import kerastuner as kt
5 from kerastuner import tuners
6
7 def MyHyperModel(hp):
8
9     image_size = (224, 224)
10
11     featurizer_mobilenet.trainable = False
12     model_input_mobilenet = featurizer_mobilenet.input
13     x = Flatten()(featurizer_mobilenet.output)
14     for i in range(hp.Int('num_layers_rnn', 2, 3)):
15         x = Dense(hp.Int('units', min_value=64, max_value=128, step=16), activation='relu')(x)
16         #x = Dropout(0.2)(x)
17     x = Dense(42, activation='softmax')(x)
18
19     model_mobilenet = Model(model_input_mobilenet, x)
20     #model_mobilenet.summary()
21     model_mobilenet.compile(loss='categorical_crossentropy',
22                             optimizer=hp.Choice('optimizer', values= ['Adam', 'RMSprop']),
23                             metrics=['acc', tf.keras.metrics.TopKCategoricalAccuracy])
24
25     return model_mobilenet
26
27 tuner = kt.Hyperband(MyHyperModel,
28                       objective = 'acc',
29                       max_epochs = 2,
30                       factor = 3,
31                       directory = '/content/drive/My Drive/sample_3000',
32                       project_name = 'mobilenet_hyper')
33
34 tuner = kt.RandomSearch(
35     MyHyperModel,
36     objective='val_acc',
37     max_trials = 2, #Change this for more iterations
38     directory='/content/drive/My Drive/sample_3000',
39     project_name='mobilenet_hyper3'
40 )
41
42 image_size = (240,240)
43 train_datagen = ImageDataGenerator(
44     rotation_range=45.0,
45     horizontal_flip=True,
46     shear_range=30.0,
47     zoom_range=0.5,
48     preprocessing_function=preprocess_input)
49
50 X_train = train_datagen.flow_from_directory('/content/drive/My Drive/sample_3000')
51

```

*← the version later on dropout tuning, and can supercede this version*

```

52 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
53 X_val = val_datagen.flow_from_directory(*content/drive/My Drive/sample_3000/
54 X_val_batch, y_val_batch = next(X_val)
55
56 logdir_mobilenet = os.path.join("/content/drive/My Drive/sample_3000/logs/mob
57 tensorboard_callback_mobilenet = tf.keras.callbacks.TensorBoard(logdir_mobile
58
59 tuner.search(X_train, epochs=1, validation_data=X_val, callbacks=[tensorboard

```

```

Found 2532 images belonging to 42 classes.
Found 630 images belonging to 42 classes.
1/80 [.....] - ETA: 0s - loss: 4.3508 - acc: 0.00
Instructions for updating:
use `tf.profiler.experimental.stop` instead.
80/80 [=====] - 2057s 26s/step - loss: 4.0510 - ac

```

## Trial complete

### Trial summary

I-Trial ID: 57634774b2d83a6cbcd3ec088809258b

I-Score: 0.03650793805718422

I-Best step: 0

### Hyperparameters:

I-num\_layers\_rnn: 2

I-optimizer: Adadelata

I-units: 80

80/80 [=====] - 62s 781ms/step - loss: 4.2620 - ac

## Trial complete

### Trial summary

I-Trial ID: df3ffeb1bd27ff63f09e911f785ef980

I-Score: 0.1190476194024086

I-Best step: 0


### Hyperparameters:

I-num\_layers\_rnn: 2

I-optimizer: Adam

I-units: 80

INFO:tensorflow:Oracle triggered exit


## ▼ --- Tensorboard - hyperparameter search ---





```

1 #mobilenet featurizer + kerastuner + tensorboard visualisation setup
2 #Get a sense of which hyperparameters make training accuracy faster
3 !pip install keras-tuner
4 import kerastuner as kt
5 from kerastuner import tuners
6 from tensorboard.plugins.hparams import api as hpp
7

```

```

8 #HP_NUM_LAYERS = hpp.HParam('num_layers', hpp.Discrete([2,3]))
9 HP_NUM_UNITS1 = hpp.HParam('num_units1', hpp.Discrete([96,128])) #CHANGE THIS
10 HP_NUM_UNITS2 = hpp.HParam('num_units2', hpp.Discrete([64,96])) #CHANGE THIS
11 HP_DROPOUT = hpp.HParam('dropout', hpp.RealInterval(0.1, 0.3))
12 HP_OPTIMIZER = hpp.HParam('optimizer', hpp.Discrete(['adam', 'sgd']))
13 METRIC_ACCURACY = 'accuracy'
14
15 with tf.summary.create_file_writer('logs4/hparam_tuning').as_default():
16     hpp.hparams_config(
17         hparams=[HP_NUM_UNITS1, HP_NUM_UNITS2, HP_DROPOUT, HP_OPTIMIZER],
18         metrics=[hpp.Metric(METRIC_ACCURACY)]
19     )
20
21 image_size = (240,240)
22 train_datagen = ImageDataGenerator(
23     rotation_range=45.0,
24     horizontal_flip=True,
25     shear_range=30.0,
26     zoom_range=0.5,
27     preprocessing_function=preprocess_input)
28
29 X_train = train_datagen.flow_from_directory('/content/drive/My Drive/sample_3')
30 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
31 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/sample_3000/')
32 X_val_batch, y_val_batch = next(X_val)
33
34 def MyHyperModel(hp):
35
36     featurizer_mobilenet.trainable = False
37     model_input_mobilenet = featurizer_mobilenet.input
38     x = Flatten()(featurizer_mobilenet.output)
39     #for i in range(hparams[HP_NUM_LAYERS]):
40     x = Dense(hparams[HP_NUM_UNITS1], activation='relu')(x)
41     x = Dense(hparams[HP_NUM_UNITS2], activation='relu')(x)
42     x = Dropout(hparams[HP_DROPOUT])(x)
43     x = Dense(42, activation='softmax')(x)
44
45     model_mobilenet = Model(model_input_mobilenet, x)
46     #model_mobilenet.summary()
47     model_mobilenet.compile(loss='categorical_crossentropy',
48                             optimizer=hparams[HP_OPTIMIZER],
49                             metrics=['accuracy', tf.keras.metrics.TopKCategoryc
50 #####CHANGE HOW MANY EPOCHS FOR EACH MODEL RUN
51     model_mobilenet.fit(
52         X_train, epochs = 10,
53         callbacks=[
54             tf.keras.callbacks.TensorBoard(logdir), # log metrics
55             hpp.KerasCallback(logdir, hparams), # log hparams
56         ],
57     )
58
59     print(model_mobilenet.evaluate(X_val))

```

hope  
this  
one used  
and  
not  
the  
above  
Hypermodel  
?

```

60 a, validationaccuracy, d = model_mobilenet.evaluate(X_val)
61
62 return validationaccuracy
63
64 def run(run_dir, hparams):
65     with tf.summary.create_file_writer(run_dir).as_default():
66         hpp.hparams(hparams)
67         validationaccuracy = MyHyperModel(hparams)
68         tf.summary.scalar(METRIC_ACCURACY, validationaccuracy, step=1)
69
70 """tuner = kt.Hyperband(MyHyperModel,
71                         objective = 'acc',
72                         max_epochs = 2,
73                         factor = 3,
74                         directory = '/content/drive/My Drive/sample_3000',
75                         project_name = 'mobilenet_hyper')"""
76
77 """tuner = kt.RandomSearch(
78     MyHyperModel,
79     objective='val_acc',
80     max_trials = 2,
81     directory='/content/drive/My Drive/sample_3000',
82     project_name='mobilenet_hyper3'
83 )"""
84
85 logdir = os.path.join("/content/drive/My Drive/sample_3000/logs/metrics", dat
86
87 #logdir_mobilenet = os.path.join("/content/drive/My Drive/sample_3000/logs/mo
88 #tensorboard_callback_mobilenet = tf.keras.callbacks.TensorBoard(logdir_mobil
89
90 #tuner.search(X_train, epochs=1, validation_data = X_val, callbacks=[tensorbo
91
92 session_num = 0
93
94 #for num_layers in HP_NUM_LAYERS.domain.values:
95 for num_units1 in HP_NUM_UNITS1.domain.values:
96     for num_units2 in HP_NUM_UNITS2.domain.values:
97         for dropout_rate in (HP_DROPOUT.domain.min_value, HP_DROPOUT.domain.max_v
98             for optimizer in HP_OPTIMIZER.domain.values:
99                 hparams = {
100                     HP_NUM_UNITS1: num_units1,
101                     HP_NUM_UNITS2: num_units2,
102                     HP_DROPOUT: dropout_rate,
103                     HP_OPTIMIZER: optimizer,
104                 }
105                 run_name = "run-%d" % session_num
106                 print('--- Starting trial: %s' % run_name)
107                 print({h.name: hparams[h] for h in hparams})
108                 run('logs4/hparam_tuning/' + run_name, hparams)
109                 session_num += 1

```

↑  
↙ can make these constants

Requirement already satisfied: keras-tuner in /usr/local/lib/python3.6/dist



```

Requirement already satisfied: keras in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: colorama in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: tqdm in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: terminaltables in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: future in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: tabulate in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: joblib<=0.11 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.6/dist-packages
Found 2532 images belonging to 42 classes.
Found 630 images belonging to 42 classes.
--- Starting trial: run-0
{'num_units1': 96, 'num_units2': 64, 'dropout': 0.1, 'optimizer': 'adam'}
Epoch 1/10
80/80 [=====] - 51s 632ms/step - loss: 4.0186 - acc: 0.0000
Epoch 2/10
80/80 [=====] - 49s 618ms/step - loss: 3.6537 - acc: 0.0000
Epoch 3/10
80/80 [=====] - 50s 625ms/step - loss: 3.5845 - acc: 0.0000
Epoch 4/10
80/80 [=====] - 50s 627ms/step - loss: 3.4908 - acc: 0.0000
Epoch 5/10
80/80 [=====] - 49s 614ms/step - loss: 3.3367 - acc: 0.0000
Epoch 6/10
80/80 [=====] - 49s 608ms/step - loss: 3.3002 - acc: 0.0000
Epoch 7/10
80/80 [=====] - 49s 609ms/step - loss: 3.1785 - acc: 0.0000
Epoch 8/10
80/80 [=====] - 49s 607ms/step - loss: 3.0675 - acc: 0.0000
Epoch 9/10
80/80 [=====] - 50s 623ms/step - loss: 3.0706 - acc: 0.0000
Epoch 10/10
80/80 [=====] - 50s 627ms/step - loss: 2.9912 - acc: 0.0000
20/20 [=====] - 6s 295ms/step - loss: 2.9766 - acc: 0.5142857432365417
[2.976576566696167, 0.17301587760448456, 0.5142857432365417]
20/20 [=====] - 6s 285ms/step - loss: 2.9766 - acc: 0.5142857432365417
--- Starting trial: run-1
{'num_units1': 96, 'num_units2': 64, 'dropout': 0.1, 'optimizer': 'sgd'}
Epoch 1/10
80/80 [=====] - 50s 630ms/step - loss: 3.6143 - acc: 0.0000
Epoch 2/10
80/80 [=====] - 51s 638ms/step - loss: 3.0925 - acc: 0.0000
Epoch 3/10
80/80 [=====] - 50s 629ms/step - loss: 2.7200 - acc: 0.0000
Epoch 4/10
80/80 [=====] - 50s 625ms/step - loss: 2.4918 - acc: 0.0000
Epoch 5/10
80/80 [=====] - 50s 628ms/step - loss: 2.2647 - acc: 0.0000
Epoch 6/10
80/80 [=====] - 50s 630ms/step - loss: 2.1491 - acc: 0.0000
Epoch 7/10
80/80 [=====] - 50s 624ms/step - loss: 1.9777 - acc: 0.0000

```



Epoch 8/10

00/00 [-----] 1 40s 617ms/step loss: 1.8620 acc: 0.0000

```
1 #Run the below individually to view and understand input shapes
2 X_train_batch, y_train_batch = next(X_train)
3
4 X_train_batch.shape
5 y_train_batch.shape
6 X_train_batch[0].shape
7 X_train_batch[0]
8 y_train_batch.shape
9 y_train_batch
10 features_batch = featurizer.predict(X_train_batch)
11 features_batch
12 features_batch[0]
13 features_batch[0].shape
```

(8, 224, 224, 3)

```
1 #Sorting by descending accuracy, we see that best hyperparameters that can be
2 #All the optimizer sgd models worked better than using adam
3 %load_ext tensorboard
4 %tensorboard --logdir logs4/hparam_tuning
```

Note: This picture shows the results from running the above cell. The above is not showing output now as the file was saved only in a temporary storage.

[Click for larger image] (<https://drive.google.com/uc?id=1rvR11t28QqWNO5Y-hiaSzKLglU7A37M4>)



```
1 %tensorboard --logdir "/content/drive/My Drive/sample_3000/logs"
```

```
1 #inception featurizer + kerastuner + tensorboard visualisation setup
2 #Get a sense of which hyperparameters make training accuracy faster
3 #!pip install keras-tuner
4 import kerastuner as kt
5 from kerastuner import tuners
6 from tensorboard.plugins.hparams import api as hpp
7
8 #HP_NUM_LAYERS = hpp.HParam('num_layers', hpp.Discrete([2,3]))
9 HP_NUM_UNITS1 = hpp.HParam('num_units1', hpp.Discrete([96, 128])) #CHANGE THIS
10 HP_NUM_UNITS2 = hpp.HParam('num_units2', hpp.Discrete([64, 96])) #CHANGE THIS
11 HP_DROPOUT = hpp.HParam('dropout', hpp.RealInterval(0.1, 0.3))
```

```

13 HP_OPTIMIZER=hpp.HParam('optimizer', hpp.Discrete(['adam', 'sgd']))
14
15 with tf.summary.create_file_writer('logs5/hparam_tuning').as_default():
16     hpp.hparams_config(
17         hparams=[HP_NUM_UNITS1, HP_NUM_UNITS2, HP_DROPOUT, HP_OPTIMIZER],
18         metrics=[hpp.Metric(METRIC_ACCURACY)]
19     )
20
21 image_size = (299,299)
22 train_datagen = ImageDataGenerator(
23     rotation_range=45.0,
24     horizontal_flip=True,
25     shear_range=30.0,
26     zoom_range=0.5,
27     preprocessing_function=preprocess_input)
28
29 X_train = train_datagen.flow_from_directory('/content/drive/My Drive/sample_3
30 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
31 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/sample_3000/
32 X_val_batch, y_val_batch = next(X_val)
33
34 def MyHyperModel(hp):
35
36     featurizer_inception.trainable = False
37     model_input_inception = featurizer_inception.input
38     x = Flatten()(featurizer_inception.output)
39     #for i in range(hparams[HP_NUM_LAYERS]):
40     x = Dense(hparams[HP_NUM_UNITS1], activation='relu')(x)
41     x = Dense(hparams[HP_NUM_UNITS2], activation='relu')(x)
42     x = Dropout(hparams[HP_DROPOUT])(x)
43     x = Dense(42, activation='softmax')(x)
44
45     model_inception = Model(model_input_inception, x)
46     #model_inception.summary()
47     model_inception.compile(loss='categorical_crossentropy',
48                             optimizer=hparams[HP_OPTIMIZER],
49                             metrics=['accuracy', tf.keras.metrics.TopKCategory
50 #####CHANGE HOW MANY EPOCHS FOR EACH MODEL RUN
51     model_inception.fit(
52         X_train, epochs = 10,
53         callbacks=[
54             tf.keras.callbacks.TensorBoard(logdir), # log metrics
55             hpp.KerasCallback(logdir, hparams), # log hparams
56         ],
57     )
58
59     print(model_inception.evaluate(X_val))
60     a, validationaccuracy, d = model_inception.evaluate(X_val)
61
62     return validationaccuracy
63

```

```

64 def run(run_dir, hparams):
65     with tf.summary.create_file_writer(run_dir).as_default():
66         hpp.hparams(hparams)
67         validationaccuracy = MyHyperModel(hparams)
68         tf.summary.scalar(METRIC_ACCURACY, validationaccuracy, step=1)
69
70     """tuner = kt.Hyperband(MyHyperModel,
71                             objective = 'acc',
72                             max_epochs = 2,
73                             factor = 3,
74                             directory = '/content/drive/My Drive/sample_3000',
75                             project_name = 'inception_hyper')"""
76
77     """tuner = kt.RandomSearch(
78         MyHyperModel,
79         objective='val_acc',
80         max_trials = 2,
81         directory='/content/drive/My Drive/sample_3000',
82         project_name='inception_hyper3'
83     )"""
84
85 logdir = os.path.join("/content/drive/My Drive/sample_3000/logs5/metrics", da
86
87 #logdir_inception = os.path.join("/content/drive/My Drive/sample_3000/logs/in
88 #tensorboard_callback_inception = tf.keras.callbacks.TensorBoard(logdir_incep
89
90 #tuner.search(X_train, epochs=1, validation_data = X_val, callbacks=[tensorbo
91
92 session_num = 0
93
94 #for num_layers in HP_NUM_LAYERS.domain.values:
95 for num_units1 in HP_NUM_UNITS1.domain.values:
96     for num_units2 in HP_NUM_UNITS2.domain.values:
97         for dropout_rate in (HP_DROPOUT.domain.min_value, HP_DROPOUT.domain.max_v
98             for optimizer in HP_OPTIMIZER.domain.values:
99                 hparams = {
100                     HP_NUM_UNITS1: num_units1,
101                     HP_NUM_UNITS2: num_units2,
102                     HP_DROPOUT: dropout_rate,
103                     HP_OPTIMIZER: optimizer,
104                 }
105                 run_name = "run-%d" % session_num
106                 print('--- Starting trial: %s' % run_name)
107                 print({h.name: hparams[h] for h in hparams})
108                 run('logs5/hparam_tuning/' + run_name, hparams)
109                 session_num += 1

```

Found 2532 images belonging to 42 classes.

Found 630 images belonging to 42 classes.

--- Starting trial: run-0

{'num\_units1': 96, 'num\_units2': 64, 'dropout': 0.1, 'optimizer': 'adam'}

Epoch 1/10

2/2000

1 - ETA: 11s - loss: 6.0820 - accuracy

```
2/80 [=====] - 67s 841ms/step - loss: 5.0029 - accuracy
80/80 [=====] - 69s 867ms/step - loss: 5.2451 - ac
Epoch 2/10
80/80 [=====] - 68s 856ms/step - loss: 3.5861 - ac
Epoch 3/10
80/80 [=====] - 67s 840ms/step - loss: 3.4501 - ac
Epoch 4/10
80/80 [=====] - 68s 845ms/step - loss: 3.3192 - ac
Epoch 5/10
80/80 [=====] - 67s 839ms/step - loss: 3.1631 - ac
Epoch 6/10
80/80 [=====] - 67s 836ms/step - loss: 3.2075 - ac
Epoch 7/10
80/80 [=====] - 67s 837ms/step - loss: 3.0683 - ac
Epoch 8/10
80/80 [=====] - 67s 834ms/step - loss: 2.9981 - ac
Epoch 9/10
80/80 [=====] - 67s 831ms/step - loss: 2.8827 - ac
Epoch 10/10
80/80 [=====] - 66s 821ms/step - loss: 2.8202 - ac
20/20 [=====] - 6s 312ms/step - loss: 2.6579 - acc
[2.6578850746154785, 0.2730158865451813, 0.5476190447807312]
20/20 [=====] - 6s 294ms/step - loss: 2.6579 - acc
--- Starting trial: run-1
{'num_units1': 96, 'num_units2': 64, 'dropout': 0.1, 'optimizer': 'sgd'}
Epoch 1/10
80/80 [=====] - 66s 831ms/step - loss: 3.8136 - ac
Epoch 2/10
80/80 [=====] - 66s 830ms/step - loss: 3.4532 - ac
Epoch 3/10
80/80 [=====] - 66s 827ms/step - loss: 3.0611 - ac
Epoch 4/10
80/80 [=====] - 66s 828ms/step - loss: 2.7727 - ac
Epoch 5/10
80/80 [=====] - 67s 833ms/step - loss: 2.4935 - ac
Epoch 6/10
80/80 [=====] - 67s 835ms/step - loss: 2.4025 - ac
Epoch 7/10
80/80 [=====] - 66s 831ms/step - loss: 2.1892 - ac
Epoch 8/10
80/80 [=====] - 66s 831ms/step - loss: 2.1046 - ac
Epoch 9/10
80/80 [=====] - 67s 833ms/step - loss: 2.0000 - ac
Epoch 10/10
80/80 [=====] - 68s 847ms/step - loss: 1.8843 - ac
20/20 [=====] - 6s 302ms/step - loss: 1.9215 - acc
[1.9214684963226318, 0.5111111402511597, 0.7714285850524902]
20/20 [=====] - 6s 299ms/step - loss: 1.9215 - acc
--- Starting trial: run-2
{'num_units1': 96, 'num_units2': 64, 'dropout': 0.3, 'optimizer': 'adam'}
Epoch 1/10
80/80 [=====] - 68s 851ms/step - loss: 4.8350 - ac
Epoch 2/10
80/80 [=====] - 67s 842ms/step - loss: 3.7359 - ac
Epoch 3/10
```

```
1 #%load_ext tensorboard
2 %tensorboard --logdir logs5/hparam_tuning
```

Note: This picture shows the results from running the above cell. The above is not showing output now as the file was saved only in a temporary storage.

[Click for larger image] (<https://drive.google.com/uc?id=1rvR11t28QgWNO5Y-hiaSzKLglU7A37M4>)



```
1 %tensorboard --logdir "/content/drive/My Drive/sample_3000/logs5"
```

### ▼ --- 3000 sample train on best model + hyperparameters ---

```
1 #Based on the results for mobilenet and inception based on the above tensorbo
2 #We will train a model using INCEPTION: We perform training on more epochs, u
3 #It did not make it to the full 200 epochs, only up to epoch 140
4
5 #adam = tf.keras.optimizers.Adam(learning_rate=1.0)
6
7 #Step 1: Load existing saved model or if there is not, create a new model
8 image_size = (299, 299)
9 model_file_inception = '/content/drive/My Drive/sample_3000/featurizer_incept
10 if os.path.isfile(model_file_inception):
11     model_inception = load_model(model_file_inception)
12 else:
13     featurizer_inception.trainable = False
14     model_input_inception = featurizer_inception.input
15     x = Flatten()(featurizer_inception.output)
16
17     x = Dense(128)(x)
18     x = LeakyReLU()(x)
19     x = Dense(96)(x)
20     x = LeakyReLU()(x)
21     x = Dropout(0.1)(x)
22     x = Dense(42, activation='softmax')(x)
23
24     model_inception = Model(model_input_inception, x)
25     model_inception.summary()
26     model_inception.compile(loss='categorical_crossentropy', optimizer='sgd', m
27
28 #Step 2: ImageDataGenerator and defining X_train, X_val
```

*Hyper version didn't include LeakyReLU...*

```

29 batch_size = 64
30 train_datagen = ImageDataGenerator(
31     rotation_range=45.0,
32     horizontal_flip=True,
33     shear_range=30.0,
34     zoom_range=0.5,
35     preprocessing_function=preprocess_input)
36 X_train = train_datagen.flow_from_directory('/content/drive/My Drive/sample_3
37 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
38 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/sample_3000/
39 X_val_batch, y_val_batch = next(X_val)
40
41 #Step 3: Run the model
42 mc_inception = ModelCheckpoint('/content/drive/My Drive/sample_3000/featurize
43 logdir_inception = os.path.join("/content/drive/My Drive/sample_3000/final_lo
44 tensorboard_callback_inception = tf.keras.callbacks.TensorBoard(logdir_incept
45 history_inception = model_inception.fit(X_train, epochs=200,
46                                     validation_data=(X_val_batch, y_val_b
47                                     callbacks=[mc_inception, tensorboard_

```

Model: "functional\_75"

Layer (type)	Output Shape	Param #	Connected
input_4 (InputLayer)	[(None, 299, 299, 3)]	0	
conv2d_94 (Conv2D)	(None, 149, 149, 32)	864	input_4[0]
batch_normalization_94 (Batch Normalization)	(None, 149, 149, 32)	96	conv2d_94[0]
activation_94 (Activation)	(None, 149, 149, 32)	0	batch_normalization_94[0]
conv2d_95 (Conv2D)	(None, 147, 147, 32)	9216	activation_94[0]
batch_normalization_95 (Batch Normalization)	(None, 147, 147, 32)	96	conv2d_95[0]
activation_95 (Activation)	(None, 147, 147, 32)	0	batch_normalization_95[0]
conv2d_96 (Conv2D)	(None, 147, 147, 64)	18432	activation_95[0]
batch_normalization_96 (Batch Normalization)	(None, 147, 147, 64)	192	conv2d_96[0]
activation_96 (Activation)	(None, 147, 147, 64)	0	batch_normalization_96[0]
max_pooling2d_4 (MaxPooling2D)	(None, 73, 73, 64)	0	activation_96[0]
conv2d_97 (Conv2D)	(None, 73, 73, 80)	5120	max_pooling2d_4[0]
batch_normalization_97 (Batch Normalization)	(None, 73, 73, 80)	240	conv2d_97[0]
activation_97 (Activation)	(None, 73, 73, 80)	0	batch_normalization_97[0]
conv2d_98 (Conv2D)	(None, 71, 71, 192)	138240	activation_97[0]
batch_normalization_98 (Batch Normalization)	(None, 71, 71, 192)	576	conv2d_98[0]

activation_98 (Activation)	(None, 71, 71, 192)	0	batch_norm
max_pooling2d_5 (MaxPooling2D)	(None, 35, 35, 192)	0	activation
conv2d_102 (Conv2D)	(None, 35, 35, 64)	12288	max_poolin
batch_normalization_102 (BatchN	(None, 35, 35, 64)	192	conv2d_102
activation_102 (Activation)	(None, 35, 35, 64)	0	batch_norm
conv2d_100 (Conv2D)	(None, 35, 35, 48)	9216	max_poolin
conv2d_103 (Conv2D)	(None, 35, 35, 96)	55296	activation
batch_normalization_100 (BatchN	(None, 35, 35, 48)	144	conv2d_100
batch_normalization_103 (BatchN	(None, 35, 35, 96)	288	conv2d_103
activation_100 (Activation)	(None, 35, 35, 48)	0	batch_norm
activation_103 (Activation)	(None, 35, 35, 96)	0	batch_norm
average_pooling2d_9 (AveragePoo	(None, 35, 35, 192)	0	max_poolin

```

1 #We notice that after the first 15 epochs, there is divergence between the tr
2 %load_ext tensorboard
3 %tensorboard --logdir "/content/drive/My Drive/sample_3000/final_logs"

```

```

1 #Run classification_report on the best model
2 batch_size = 630
3 image_size = (299, 299)
4
5 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
6 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/sample_3000/'
7 X_val_batch, y_val_batch = next(X_val)
8
9 best_model = load_model('/content/drive/My Drive/sample_3000/featurizer_incep
10 pred_test = best_model.predict(X_val_batch)
11 print(classification_report(y_val_batch.argmax(axis=1), pred_test.argmax(axis:

```

Found 630 images belonging to 42 classes.

	precision	recall	f1-score	support
0	0.65	0.69	0.67	16
1	0.54	0.88	0.67	16
2	0.62	0.50	0.55	16
3	0.60	0.38	0.46	16
4	0.76	0.81	0.79	16
5	0.67	0.88	0.76	16
6	0.58	0.44	0.50	16
7	0.50	0.69	0.58	16
8	0.65	0.69	0.67	16
9	0.60	0.75	0.67	16
10	0.64	0.44	0.52	16
11	1.00	0.55	0.71	11
12	0.41	0.88	0.56	16
13	0.82	0.88	0.85	16
14	0.91	0.62	0.74	16
15	0.73	0.50	0.59	16
16	0.78	0.44	0.56	16
17	0.73	0.89	0.80	9
18	0.29	0.38	0.33	13
19	0.57	0.25	0.35	16
20	0.54	0.44	0.48	16
21	0.67	0.75	0.71	16
22	0.79	0.69	0.73	16
23	0.62	0.33	0.43	15
24	0.80	0.75	0.77	16
25	0.92	0.69	0.79	16
26	0.31	0.69	0.42	16
27	0.48	0.81	0.60	16
28	0.70	0.93	0.80	15
29	0.80	0.62	0.70	13
30	0.71	0.94	0.81	16
31	0.35	0.75	0.48	16
32	1.00	0.23	0.38	13
33	1.00	0.33	0.50	3
34	0.58	0.88	0.70	16
35	0.88	0.44	0.58	16
36	0.25	0.12	0.17	16
37	0.75	0.60	0.67	10
38	0.80	0.25	0.38	16
39	0.38	0.50	0.43	16
40	0.75	0.38	0.50	16
41	0.43	0.19	0.26	16
accuracy			0.60	630
macro avg	0.66	0.59	0.59	630
weighted avg	0.64	0.60	0.59	630

### ▼ --- Finetuning featurizer ---

1 #Finetuning – Inspect which layers to finetune



```
2 featurizer_inception = InceptionV3(input_shape = (299, 299, 3), include_top =
3 featurizer_inception.summary()
```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/87916544/87910968> [=====] - 2s 0us/step  
Model: "inception\_v3"

Layer (type)	Output Shape	Param #	Connected
input_1 (InputLayer)	[(None, 299, 299, 3)]	0	
conv2d (Conv2D)	(None, 149, 149, 32)	864	input_1[0]
batch_normalization (Batch Normalization)	(None, 149, 149, 32)	96	conv2d[0]
activation (Activation)	(None, 149, 149, 32)	0	batch_normalization
conv2d_1 (Conv2D)	(None, 147, 147, 32)	9216	activation
batch_normalization_1 (Batch Normalization)	(None, 147, 147, 32)	96	conv2d_1[0]
activation_1 (Activation)	(None, 147, 147, 32)	0	batch_normalization_1
conv2d_2 (Conv2D)	(None, 147, 147, 64)	18432	activation_1
batch_normalization_2 (Batch Normalization)	(None, 147, 147, 64)	192	conv2d_2[0]
activation_2 (Activation)	(None, 147, 147, 64)	0	batch_normalization_2
max_pooling2d (MaxPooling2D)	(None, 73, 73, 64)	0	activation_2
conv2d_3 (Conv2D)	(None, 73, 73, 80)	5120	max_pooling2d
batch_normalization_3 (Batch Normalization)	(None, 73, 73, 80)	240	conv2d_3[0]
activation_3 (Activation)	(None, 73, 73, 80)	0	batch_normalization_3
conv2d_4 (Conv2D)	(None, 71, 71, 192)	138240	activation_3
batch_normalization_4 (Batch Normalization)	(None, 71, 71, 192)	576	conv2d_4[0]
activation_4 (Activation)	(None, 71, 71, 192)	0	batch_normalization_4
max_pooling2d_1 (MaxPooling2D)	(None, 35, 35, 192)	0	activation_4
conv2d_8 (Conv2D)	(None, 35, 35, 64)	12288	max_pooling2d_1
batch_normalization_8 (Batch Normalization)	(None, 35, 35, 64)	192	conv2d_8[0]
activation_8 (Activation)	(None, 35, 35, 64)	0	batch_normalization_8
conv2d_6 (Conv2D)	(None, 35, 35, 48)	9216	activation_8
conv2d_9 (Conv2D)	(None, 35, 35, 96)	55296	conv2d_6
batch_normalization_6 (Batch Normalization)	(None, 35, 35, 48)	144	conv2d_9[0]
activation_6 (Activation)	(None, 35, 35, 48)	0	batch_normalization_6

batch_normalization_9 (Batch Normalization)	(None, 35, 35, 96)	200	conv2d_9 [0]
activation_6 (Activation)	(None, 35, 35, 48)	0	batch_norm
activation_9 (Activation)	(None, 35, 35, 96)	0	batch_norm

```

1 #Unfreeze all models after "mixed6"
2 featurizer_inception.trainable = False
3 unfreeze = False
4 for l in featurizer_inception.layers:
5     print(l.name + ' unfrozen: ' + str(unfreeze))
6     if unfreeze:
7         l.trainable = True
8     if l.name == 'mixed6':
9         unfreeze = True

```

```

input_1 unfrozen: False
conv2d unfrozen: False
batch_normalization unfrozen: False
activation unfrozen: False
conv2d_1 unfrozen: False
batch_normalization_1 unfrozen: False
activation_1 unfrozen: False
conv2d_2 unfrozen: False
batch_normalization_2 unfrozen: False
activation_2 unfrozen: False
max_pooling2d unfrozen: False
conv2d_3 unfrozen: False
batch_normalization_3 unfrozen: False
activation_3 unfrozen: False
conv2d_4 unfrozen: False
batch_normalization_4 unfrozen: False
activation_4 unfrozen: False
max_pooling2d_1 unfrozen: False
conv2d_8 unfrozen: False
batch_normalization_8 unfrozen: False
activation_8 unfrozen: False
conv2d_6 unfrozen: False
conv2d_9 unfrozen: False
batch_normalization_6 unfrozen: False
batch_normalization_9 unfrozen: False
activation_6 unfrozen: False
activation_9 unfrozen: False
average_pooling2d unfrozen: False
conv2d_5 unfrozen: False
conv2d_7 unfrozen: False
conv2d_10 unfrozen: False
conv2d_11 unfrozen: False
batch_normalization_5 unfrozen: False
batch_normalization_7 unfrozen: False
batch_normalization_10 unfrozen: False
batch_normalization_11 unfrozen: False
activation_5 unfrozen: False
activation_7 unfrozen: False
activation_10 unfrozen: False
activation_11 unfrozen: False

```

```

mixed0 unfrozen: False
conv2d_15 unfrozen: False
batch_normalization_15 unfrozen: False
activation_15 unfrozen: False
conv2d_13 unfrozen: False
conv2d_16 unfrozen: False
batch_normalization_13 unfrozen: False
batch_normalization_16 unfrozen: False
activation_13 unfrozen: False
activation_16 unfrozen: False
average_pooling2d_1 unfrozen: False
conv2d_12 unfrozen: False
conv2d_14 unfrozen: False
conv2d_17 unfrozen: False
conv2d_18 unfrozen: False
batch_normalization_12 unfrozen: False
batch_normalization_14 unfrozen: False
batch_normalization_17 unfrozen: False
batch_normalization_18 unfrozen: False
activation_12 unfrozen: False

```

```

1 #Finetuning for 10 epochs at a low learning rate
2 #Somehow, the tensorboard command produced an error, so we will comment it out
3 from tensorflow.keras.optimizers import SGD
4
5 model_file_inception = '/content/drive/My Drive/sample_3000/featurizer_inception.h5'
6 model_inception = load_model(model_file_inception)
7
8 unfreeze = False
9 for l in featurizer_inception.layers:
10     if unfreeze:
11         l.trainable = True
12     if l.name == 'mixed6':
13         unfreeze = True
14
15 model_input_inception = featurizer_inception.input
16 x = Flatten()(featurizer_inception.output)
17
18 x = Dense(128)(x)
19 x = LeakyReLU()(x)
20 x = Dense(96)(x)
21 x = LeakyReLU()(x)
22 x = Dropout(0.1)(x)
23 x = Dense(42, activation='softmax')(x)
24
25 model_inception = Model(model_input_inception, x)
26 model_inception.compile(loss='categorical_crossentropy', optimizer=SGD(learning_rate=0.001))
27
28 batch_size = 64
29 image_size = (299, 299)
30 train_datagen = ImageDataGenerator(
31     rotation_range=45.0,
32     horizontal_flip=True,
33     shear_range=30.0,

```

```

34     zoom_range=0.5,
35     preprocessing_function=preprocess_input)
36 X_train = train_datagen.flow_from_directory('/content/drive/My Drive/sample_3
37 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
38 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/sample_3000/
39 X_val_batch, y_val_batch = next(X_val)
40
41 mc = ModelCheckpoint('/content/drive/My Drive/sample_3000/featurizer_inceptio
42 #logdir_inception = os.path.join("/content/drive/My Drive/sample_3000/final_l
43 #tensorboard_callback_inception = tf.keras.callbacks.TensorBoard(logdir_incep
44 #,tensorboard_callback_inception
45 history = model_inception.fit(X_train, epochs=10,
46                             validation_data=(X_val_batch, y_val_batch),
47                             callbacks=[mc])

```

Found 2532 images belonging to 42 classes.

Found 630 images belonging to 42 classes.

Epoch 1/10

40/40 [=====] - ETA: 0s - loss: 0.9954 - acc: 0.70

40/40 [=====] - 423s 11s/step - loss: 0.9954 - acc:

Epoch 2/10

40/40 [=====] - 64s 2s/step - loss: 0.9879 - acc:

Epoch 3/10

40/40 [=====] - 64s 2s/step - loss: 0.9070 - acc:

Epoch 4/10

40/40 [=====] - 64s 2s/step - loss: 0.8851 - acc:

Epoch 5/10

40/40 [=====] - 62s 2s/step - loss: 0.8351 - acc:

Epoch 6/10

40/40 [=====] - 62s 2s/step - loss: 0.8182 - acc:

Epoch 7/10

40/40 [=====] - 64s 2s/step - loss: 0.7958 - acc:

Epoch 8/10

40/40 [=====] - 66s 2s/step - loss: 0.7811 - acc:

Epoch 9/10

40/40 [=====] - 63s 2s/step - loss: 0.7633 - acc:


Epoch 10/10

40/40 [=====] - 62s 2s/step - loss: 0.7131 - acc:


```

1 #The below can be run if there were no time-out that occurred. But we can see
2 plt.plot(history.history['acc'], label='train set')
3 plt.plot(history.history['val_acc'], label='validation set')
4 plt.xlabel('Epochs')
5 plt.ylabel('Accuracy')
6 plt.title('Model accuracy')
7 plt.legend(['Train', 'Test'], loc='upper left')
8 plt.show()
9
10 plt.plot(history.history['loss'], label='train set')
11 plt.plot(history.history['val_loss'], label='validation set')
12 plt.xlabel('Epochs')
13 plt.ylabel('Loss')
14 plt.title('Loss')
15 plt.legend()
16 plt.show()

```



*Starting to overfit*



▼ --- Global dataset ---

```

1 #So far, we have run on the dataset of 3000 images
2 #Now we run on the global dataset of 100K images
3 #Split global dataset into train and validation sets - Creating CSV file to n
4 df_traincsv_marker = df_traincsv
5 df_traincsv_marker['selected_for_validation'] = 0
6 df_traincsv_marker.head()

```

	filename	category	selected_for_validation
0	45e2d0c97f7bdf8cbf3594beb6fdcdca0.jpg	3	0
1	f74d1a5fc2498bbbfa045c74e3cc333e.jpg	3	0
2	f6c172096818c5fab10ecae722840798.jpg	3	0
3	251ffd610399ac00fea7709c642676ee.jpg	3	0
4	73c7328b8eda399199fededec6e4badaf.jpg	3	0

```

1 #DO NOT RUN THIS CELL AFTER FIRST RUN
2 rolling = 20
3 for two_digit in folder_names[20:]:
4     path = "/content/drive/My Drive/validation/" + two_digit
5     os.mkdir(path)
6
7     temp_df = df_traincsv_marker.loc[df_traincsv_marker.category==rolling,:].sa
8     temp_index = temp_df.index
9     df_traincsv_marker.loc[temp_index, 'selected_for_validation'] = 1
10    temp_df.reset_index(inplace=True)
11
12    for i in range(len(temp_df)):
13        f = temp_df.loc[i,'filename']
14        shutil.move("/content/drive/My Drive/train/train/train/" + two_digit + '/'
15
16    rolling+=1

```

```

1 featurizer_mobilenet = MobileNetV2(input_shape=(224, 224, 3), include_top=False)
2 featurizer_inception = InceptionV3(input_shape = (299, 299, 3), include_top =

```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/9412608/9406464> [=====] - 0s 0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/87916544/87910968> [=====] - 2s 0us/step

```

1 #As the global dataset is large, we expect that the colab session will run out
2 #We create a function that can re-used after the colab session runs out
3 #RUN THIS CELL AFTER COLAB 6-HOUR RUNTIME
4 #optimizer=SGD(learning_rate=0.9)
5 adam = tf.keras.optimizers.Adam(learning_rate=1.0)
6

```

consider making a function with parameter for dataset size

you can also sample from larger set and fine tune / until

```

7 def load_and_run_100k(the_model):
8     #Step 1: Load existing saved model or if there is not, create a new model
9     if the_model == 'mobilenet':
10         image_size = (224, 224)
11         #model_file_mobilenet_savebest = '/content/drive/My Drive/featurizer_mobi
12         #model_file_mobilenet_savefreq = '/content/drive/My Drive/featurizer_mobi
13         model_file_mobilenet_savefreq = '/content/drive/My Drive/featurizer_mobil
14
15         if os.path.isfile(model_file_mobilenet_savefreq):
16             model_mobilenet = load_model(model_file_mobilenet_savefreq)
17
18         else:
19             featurizer_mobilenet.trainable = False
20             model_input_mobilenet = featurizer_mobilenet.input
21             x = Flatten()(featurizer_mobilenet.output)
22
23             x = Dense(128)(x)
24             x = LeakyReLU()(x)
25             x = Dense(64)(x)
26             x = LeakyReLU()(x)
27             x = Dense(42, activation='softmax')(x)
28
29             model_mobilenet = Model(model_input_mobilenet, x)
30             model_mobilenet.summary()
31             model_mobilenet.compile(loss='categorical_crossentropy', optimizer=adam
32
33
34     elif the_model == 'inception':
35         image_size = (299, 299)
36         #model_file_inception_savebest = '/content/drive/My Drive/featurizer_ince
37         model_file_inception_savefreq = '/content/drive/My Drive/featurizer_incep
38
39         if os.path.isfile(model_file_inception_savefreq):
40             print('yes')
41             model_inception = load_model(model_file_inception_savefreq)
42         else:
43             featurizer_inception.trainable = False
44             model_input_inception = featurizer_inception.input
45             x = Flatten()(featurizer_inception.output)
46
47             x = Dense(128)(x)
48             x = LeakyReLU()(x)
49             x = Dense(64)(x)
50             x = LeakyReLU()(x)
51             x = Dense(42, activation='softmax')(x)
52
53             model_inception = Model(model_input_inception, x)
54             model_inception.summary()
55             model_inception.compile(loss='categorical_crossentropy', optimizer=adam
56
57
58     #Step 2: ImageDataGenerator and defining X_train, X_val

```

*train accuracy goal reached.*

```

59 batch_size = 64
60
61 train_datagen = ImageDataGenerator(
62     rotation_range=45.0,
63     horizontal_flip=True,
64     shear_range=30.0,
65     zoom_range=0.5,
66     preprocessing_function=preprocess_input)
67 X_train = train_datagen.flow_from_directory('/content/drive/My Drive/train/')
68
69 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
70 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/validation')
71 X_val_batch, y_val_batch = next(X_val)
72
73 #Step 3: Run the model
74 if the_model == 'mobilenet':
75     mc_mobilenet_savebest = ModelCheckpoint('/content/drive/My Drive/featuriz')
76     mc_mobilenet_savefreq = ModelCheckpoint('/content/drive/My Drive/featuriz')
77     logdir_mobilenet = os.path.join("/content/drive/My Drive/logs/mobilenet",
78     tensorboard_callback_mobilenet = tf.keras.callbacks.TensorBoard(logdir_mo
79     history_mobilenet = model_mobilenet.fit(X_train, epochs=10,
80         validation_data=(X_val_batch, y_val_batch),
81         callbacks=[mc_mobilenet_savebest, mc_mobilenet_savefreq,
82     file_history_mobilenet = open('/content/drive/My Drive/file_history_mobil
83     pickle.dump(history_mobilenet, file_history_mobilenet)
84     file_history_mobilenet.close()
85
86 elif the_model == 'inception':
87     mc_inception_savebest = ModelCheckpoint('/content/drive/My Drive/featuriz')
88     mc_inception_savefreq = ModelCheckpoint('/content/drive/My Drive/featuriz')
89     logdir_inception = os.path.join("/content/drive/My Drive/logs/inception",
90     tensorboard_callback_inception = tf.keras.callbacks.TensorBoard(logdir_in
91     history_inception = model_inception.fit(X_train, epochs=10,
92         validation_data=(X_val_batch, y_val_batch),
93         callbacks=[mc_inception_savebest, mc_inception_savefreq,
94     file_history_inception = open('/content/drive/My Drive/file_history_incep
95     pickle.dump(history_inception, file_history_inception)
96     file_history_inception.close()

```



```
1 load_and_run_100k("mobilenet")
```

```
yes
Found 89640 images belonging to 42 classes.
Found 15742 images belonging to 42 classes.
Epoch 1/100
  1/1401 [.....] - ETA: 0s - loss: 313250.5625 -
Instructions for updating:
use `tf.profiler.experimental.stop` instead.
 377/1401 [=====>.....] - ETA: 17:11:14 - loss: 290402.4
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-5-e938c1895d8d> in <module>()
----> 1 load_and_run_100k("mobilenet")
```

```
----- 9 frames -----
/usr/local/lib/python3.6/dist-packages/tensorflow/python/eager/execute.py
in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    58     ctx.ensure_initialized()
    59     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name,
op_name,
----> 60                                     inputs, attrs, num_outputs)
    61 except core._NotOkStatusException as e:
    62     if name is not None:
```

```
KeyboardInterrupt:
```

```
1 %reload_ext tensorboard
2 %tensorboard --logdir "/content/drive/My Drive/logs"
```

```
1 #As the dataset is large and during training it did not train through 1 epoch
2 #This cell gets the validation accuracy if you specify which saved .h5 model
3 #We notice that on some classes, the model was not able to get any prediction
4 #And we note that the model did not even see one full epoch, much less get tr
5 #'/content/drive/My Drive/featurizer_mobilenet_savefreq20200929-113939-291276
6 #'/content/drive/My Drive/featurizer_inception_savefreq20200929-231315-323220
7 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
8 batch_size = int(input('Batch size: '))
9 if input('Type m for mobilenet: ') == "m":
10     image_size = (224, 224)
11 else:
12     image_size = (299, 299)
13 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/validation',
14 X_val_batch, y_val_batch = next(X_val)
15 loaded_model = input('Trained model file path: ')
16 try_model = load_model(loaded_model)
17 pred_test = try_model.predict(X_val_batch)
18 print(classification_report(y_val_batch.argmax(axis=1), pred_test.argmax(axis=
```

Batch size: 256

Type m for mobilenet: m

Found 15742 images belonging to 42 classes.

Trained model file path: /content/drive/My Drive/featurizer\_mobilenet\_savef  
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/\_classification.py:1  
\_warn\_prf(average, modifier, msg\_start, len(result))

	precision	recall	f1-score	support
0	0.00	0.00	0.00	10
1	0.54	0.88	0.67	8
2	0.40	0.22	0.29	9
3	0.25	0.12	0.17	8
4	1.00	0.83	0.91	6
5	0.50	0.43	0.46	7
6	0.67	0.25	0.36	8
7	0.75	0.60	0.67	5
8	0.25	0.33	0.29	6
9	0.00	0.00	0.00	7
10	0.40	0.57	0.47	7
11	0.00	0.00	0.00	2
12	0.15	0.60	0.24	5
13	0.50	0.86	0.63	7
14	0.67	0.67	0.67	6
15	0.50	0.43	0.46	7
16	1.00	0.20	0.33	10
17	1.00	0.12	0.22	8
18	0.00	0.00	0.00	2
19	0.33	0.50	0.40	4
20	0.00	0.00	0.00	5
21	0.00	0.00	0.00	3
22	0.40	0.67	0.50	3
23	0.38	0.60	0.46	5
24	0.57	0.57	0.57	7
25	0.50	0.50	0.50	10
26	0.00	0.00	0.00	5
27	0.00	0.00	0.00	10
28	0.67	0.67	0.67	6
29	1.00	0.67	0.80	6
30	0.50	0.38	0.43	8
31	0.00	0.00	0.00	8
32	0.40	0.57	0.47	7
33	0.12	0.50	0.20	2
34	0.25	0.67	0.36	6
35	1.00	0.54	0.70	13
36	0.18	0.67	0.29	3
37	0.25	0.25	0.25	4
38	0.00	0.00	0.00	5
39	1.00	0.33	0.50	3
40	0.00	0.00	0.00	1
41	0.00	0.00	0.00	4
accuracy			0.37	256
macro avg	0.38	0.36	0.33	256
weighted avg	0.44	0.37	0.36	256

### ▼ --- Trying model trained on 3000 samples on global validation set ---

```
1 #We try using the best model trained using the sample of 3000, using it to pr
2 #The results was surprisingly good, achieving an accuracy of 0.61
3
4 #/content/drive/My Drive/sample_3000/featurizer_inception_finetune.h5
5 val_datagen = ImageDataGenerator(preprocessing_function=preprocess_input)
6 batch_size = int(input('Batch size: '))
7 if input('Type m for mobilenet: ') == "m":
8     image_size = (224, 224)
9 else:
10    image_size = (299, 299)
11 X_val = val_datagen.flow_from_directory('/content/drive/My Drive/validation',
12 X_val_batch, y_val_batch = next(X_val)
13 loaded_model = input('Trained model file path: ')
14 try_model = load_model(loaded_model)
15 pred_test = try_model.predict(X_val_batch)
16 print(classification_report(y_val_batch.argmax(axis=1), pred_test.argmax(axis=1))
```

Batch size: 256

Type m for mobilenet: i

Found 15742 images belonging to 42 classes.

Trained model file path: /content/drive/My Drive/sample\_3000/featurizer\_inc  
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/\_classification.py:1

\_warn\_prf(average, modifier, msg\_start, len(result))

	precision	recall	f1-score	support
0	0.50	0.33	0.40	9
1	0.70	0.70	0.70	10
2	0.50	0.33	0.40	6
3	0.54	0.64	0.58	11
4	1.00	0.83	0.91	6
5	0.67	0.67	0.67	6
6	0.67	0.57	0.62	7
7	0.86	0.67	0.75	9
8	1.00	0.62	0.77	8
9	0.50	0.67	0.57	3
10	0.45	0.83	0.59	6
11	1.00	0.50	0.67	4
12	0.75	0.60	0.67	10
13	0.86	1.00	0.92	6
14	0.86	0.86	0.86	7
15	0.86	0.86	0.86	7
16	0.67	0.80	0.73	5
17	0.67	0.50	0.57	4
18	0.50	0.60	0.55	5
19	0.33	0.67	0.44	3
20	0.33	0.40	0.36	5
21	0.80	0.80	0.80	5
22	1.00	0.40	0.57	5
23	0.50	0.80	0.62	5
24	0.50	0.67	0.57	3
25	0.71	0.83	0.77	6
26	0.40	0.33	0.36	6
27	0.67	0.80	0.73	5
28	1.00	0.71	0.83	7
29	0.71	1.00	0.83	5
30	0.90	1.00	0.95	9
31	0.50	0.25	0.33	4
32	0.67	0.67	0.67	3
33	0.00	0.00	0.00	0
34	0.33	0.50	0.40	8
35	0.67	0.29	0.40	7
36	0.57	0.25	0.35	16
37	0.00	0.00	0.00	3
38	0.22	0.50	0.31	4
39	0.60	0.75	0.67	4
40	0.43	0.60	0.50	5
41	0.25	0.33	0.29	9
accuracy			0.61	256
macro avg	0.61	0.60	0.58	256
weighted avg	0.64	0.61	0.60	256

## Section 6: Evaluate Metrics



The cells above have a mixture of Section 5 and 6, as there are a few rounds of model creation and performance evaluation. Please see the above comments in green for more comments

### Preliminary testing of hyperparameters

- Worst combination (accuracy 0.025397): Inception + 96 units + 96 units + adam + dropout(0.3)
- Best combination (accuracy 0.53125): Inception + 128 units + 96 units + sgd + dropout(0.1)

### General evaluation of metrics:

- Train on 3000 sample + Inception featurizer (trainable = False) + best hyperparameters = **0.60** accuracy (on 630 images from 3000 sample validation set)
- Train on global dataset (less than 1 epoch) + Inception featurizer (trainable = False) + hyperparameters (128 units + 64 units + leaky relu + adam) = **0.37** accuracy (on 256 images from global validation set)
- Train on 3000 sample + Inception featurizer (trainable = True from mixed 7 onwards) + best hyperparameters = **0.61** accuracy (on 256 images from global validation set)

Good attempt and summary!

Some suggestions to consider

- Choose 1 version of Hypermodel and make it reflect your actual model (e.g. LeakyRelu...)
- Make the dataset logic into a function so you can easily change sizes without rewriting.
- If dataset too large to load at once, try sampling it many times & update model progressively.
- The code to throw out what is

- cleanup the not used

## ▼ Section 7: Observations and analysis

### 1. Conclusion from metrics

- As there are 42 categories for this problem / dataset, we note that an accuracy of 0.0238 ( $1/42$ ) is a random guess on which category an image belongs to. To have an accuracy of 0.61 could be considered to be pretty good in this context. ✓ *good benchmark*
- We note also that the images in the dataset are not all clean - There are some pictures that are not taken too well, and this simulates what an actual image taken by a person might be like. Per the description from kaggle, there are also some mislabelled images.
- There is some degree of imbalance in the dataset, but not wildly so. So in the sampling of data, we have ignored this degree of imbalance. ✓
- The Tensorboard hyperparameter search has been extremely useful in giving us a rough gauge on which hyperparameters would do well on our dataset. For example, all the models that had SGD as the optimizer performed better than those using adam. We can then be more confident in selecting hyperparameters to perform longer training on (more epochs). ✓

### 2. Proposed improvements

- We had decided to unfreeze inceptionv3 featurizer layers from mixed7 onwards. More tests could be run to determine the results of gradually unfreezing layers infinetuning.
- We had run training and validation on less than one epoch of the global dataset, due to shortage of time. The ideal scenario is to perform training and validation on perhaps 15 or so epochs of the global dataset.
- The model trained on 3000 samples (inceptionv3, best hyperparameters, finetuning featurizer layers from from mixed7 onwards) did surprisingly well on a batch of 256 on the global validation set. This shows, we can probably train on a bigger model of about 20,000 samples using the same methodology and probably get better results. We may not need to run training on the entire global training images we had defined, and still obtain decent results.

